

Proposals for experimental projects

Graz, Nov. and Dec. 2025



Intro experiments



- Course set-up
 - Intro+background
 - Experiments (1 or 2 persons), brief presentations
 - Short report (1 person)
 - Assessment based on active participation and report
- ASR is way too broad to be addressed in just a few weeks
 - You might choose topic yourself or from precooked list
 - Keep the research question doable
 - Keep it feasible

Python scripts available for precooked options

- In github repo [ltenbosch2/Graz_github](#)

Three (four) proposals

- Bridging phonetics and ASR/deep learning
 - Locating phonetic structure in latent representations
- Prompting Whisper
 - ASR is context dependent
- Behavior of transformer based networks on small (user-build) data sets
 - E.g., boolean mappings
 - Regular expressions: <https://aclanthology.org/2020.emnlp-main.576.pdf>
- Reasoning
 - How to validate the outcome of reasoning models?

List of experimental options

1. Bridging phonetics and ASR/deep learning

- Start by using pretrained wav2vec2.0 via Huggingface
- Script (github): [try_Wav2Vec2_latent_logit.py](#)
 - Choose dataset
 - Process mono wave file (16 kHz)
 - Extract hidden layer k , and logit tensor
 - Select two vowels (for example) that minimally differ acoustically, e.g.
 - **Bahn** [ba:n] – "track, railway" – **Bonn** [bɔn] – (city name)
 - **Beet** [be:t] – "garden bed" – **Bett** [bɛt] – "bed"
 - **fühlen** ['fy:lən] – "to feel" – **füllen** ['fylən] – "to fill"
 - **Ofen** ['o:fn] – "oven" – **offen** ['ɔfn] – "open"
 - Use logit matrix to locate the corresponding time stamps
 - Use these time stamps to build two distributions of vectors
 - Try to measure the difference (as function of k) by LDA or similar method

List of experimental options (cont'd)

2. Whisper prompting

- Present a wave file, keep this file fixed and investigate what effect text prompting has on Whisper output
- Scripts (github):
- `try_Whisper_from_OpenAI.py # OpenAI API KEY required!!`
- OR: `try_Whisper_from_HuggingFace_multiple_languages.py # this needs some further engineering to get it working`
- Follow the recipe in following link
https://cookbook.openai.com/examples/whisper_prompts_guide
- Helpful to prompt Whisper to obtain a certain transcription style
- Keep experiments doable and transparent

List of experimental options (cont'd)

3. Behavior of transformer-based networks on small (user-build) data sets

- Small data sets (training, validation, test sets)
- Minimal functional transformer models

Various proposals – take them as starting point:

- **Scripts in .txt file (github):** `Graz_regular_expressions_training_transformers.txt`
 - Data are created in `Graz_regular_expressions_create_training_test_sets.txt`
 - MA_002 (data also in github) A*B versus B*A, in random **real** vector sequence
 - MA_003, ABC versus ACB in random **integer** vector sequence (dim = 5)
 - MA_003b, same as MA_003 but with dim = 3
 - MA_004, random integer sequence → Boolean (whether the last vector is ‘between’ first two vectors, or not)
- **Graz_various_transformer_studies.txt**
 - 4 specific cases (from vector sequence to vector, to Boolean, character sequences)

List of experimental options (cont'd)

4. Use of a Large Reasoning Model (LRM)

See (github): [Reasoning_deepseek_simple.doc](#)

- Select an LRM (minimum DeepSeek-v1, install ollama)
- Select say 10 (simple to hard) questions directly related to this course (speech recognition, AI, transformers)
- Let the LRM provide a reasoning for all questions
- Assess these answers in detail
 - Define your assessment method – based on knowledge/literature
 - Provide a conclusion about the quality of the LRM chosen
- Defend your choices in model selection, method and data selection

Reasoning: example (depends on # free GB)

- install ollama
- Run any of your preferred models, e.g. deepseek (many versions)
- ollama run deepseek-coder-v2

Q: Do you see how the towers of Hanoi problem can be solved for all $N > 1$ by using recursion?

A: “Thinking... Alright, so I'm trying to understand how to solve the Tower of Hanoi problem recursively for any number of disks greater than one. From what I remember, the classic Tower of Hanoi has three rods and three disks: two large disks and a small one. The goal is to move all the disks from the starting rod to another rod while following these rules:... [etc.]”

Short report

- Report of 3 pages proceedings format (2-column format)
 - Approx 2000 words
 - Intro + RQ
 - Method
 - Experimental data + set-up
 - Results
 - References



Software?

- Enormous amount of scripts available on the web
- https://docs.pytorch.org/audio/stable/tutorials/speech_recognition_pipeline_tutorial.html
- Huggingface
- (rapidly changing – you can basically choose yourself)

Computational platforms

Up to you:

Cluster

Jupyter

SageMath

LaTeX

Cocalc:
<https://cocalc.com/>

Own laptop
• Windows, Linux

For those interested in more involving experiments

- There are (too) many options!
- E.g. Compare wav2vec2.0 **without** lexicon versus **with** lexicon and LM
 - LM of small order (e.g., 3), via KenLM toolkit
<https://kheafield.com/code/kenlm/>
 - finetuning XLSR model (Facebook, 2022) by fairseq toolkit (Ott et al., 2019)
 - Fairseq is a toolkit. With Fairseq you can train models that achieve state-of-the-art performance on machine translation and summarization tasks. It includes pre-trained models.